

Dominance-Support for the Hikita $\star$ -product: all lengths, exact up-set support

Rick

2026-10-01

Author: Rick
Date: 2026-10-01
Recipient: Clio Vega
Title: Dominance-Support for the Hikita $\star$ -product: all lengths, exact up-set support
Source: `proofs/2026-09-30-day214-DS-all-lengths-PROVED.md`
WIP commit: This proof corresponds to work-in-progress commit `9b9db51` (the proof file was last changed in `b5ccc50`, which added Theorem 2 / §8; this PDF itself is added in a follow-up commit).
Registry: `proofs/registry/hikita-star-dominance-support.json`, root node, trust: proved

Original header (Day 214 PROVE). Date: 2026-09-30 (deep-work session). Author: Rick.
Status: PROVED for every partition λ , **including the stretch goal:** the support is exactly the full up-set, and $\text{val}_s c_{\lambda\mu} = n(\mu)$ (§8). That includes DS_3 (the PROVE.md target), all longer lengths, and the operator-level version for every $e_k \star e_\mu$.
Inputs: 207b $(A_k) + (K_k)$, which give the subset formula (0.1) below, and Lemma 3.3 ($e_k \star 1 = e_k$). Everything else is textbook symmetric-function combinatorics (Macdonald I §§1, 6).
Not used: (TC), $(\star\ell)$, the Wick kernel, residues.
Scripts: `scripts/day214/` (exact sympy, all ALL OK).

Informal summary. PROVE.md was bracing for a cross- b cancellation mechanism. There isn't one to find. **Degree bounds don't care about cancellation.** Each subset term of $e_k \star F$ is a rational function whose u -degree under $x_S \rightarrow u \cdot x_S$ is at most $\deg_S F + |A \cap S|$, because every kernel factor a_{ij} has degree 0. That is the whole support proof. The (TC) termwise violations were an artefact of expanding in $E(t^i z)E(t^j w)$. The *subset* terms never violate anything. Third time the “don't reach for machinery” lesson fires; this time the lever is a degree count, not a residue.

0. Setup and statement

- $s = q^{-1}$. We work in $\Lambda_m \otimes \mathbb{Q}(s, t)$, writing $K := \mathbb{Q}(s, t)$.
- $a_{ij} := (x_i - tx_j)/(x_i - x_j)$.

- For $A \subseteq [m]$, $\Pi_A^\times := \prod_{i \in A, j \notin A} a_{ij}$ and $X_A := \prod_{a \in A} x_a$.

(0.1) Subset formula. For symmetric F and $1 \leq k \leq m$,

$$E_k F := e_k \star F = t^{-\binom{k}{2}} e_k(Y) \bullet F = \sum_{|A|=k} \Pi_A^\times \cdot X_A \cdot F(X_{A^c}, sX_A). \quad (0.1)$$

Justification.

- 207b (A_k) gives $e_k(Y)F = t^{\binom{k}{2}} \sigma^{(k)} \pi^k F$.
- $\pi^k F = x_1 \cdots x_k F(x_{k+1..m}, sx_1..sx_k)$, which is symmetric in the head and in the tail.
- 207b (K_k) gives $\sigma^{(k)} G = \sum_A G^{(A)} \Pi_A^\times$.
- This is the same formula 207b §4 and (TC) §1 start from.
- Re-checked here: `check_vs_207b.py` compares (0.1) with the 207b closed form symbolically in s, t for $k \leq 3, r \leq 3$.

Because F is symmetric, $F(X_{A^c}, sX_A)$ is simply F with $x_a \mapsto sx_a$ for $a \in A$. The order of the arguments is irrelevant.

Definition. $e_\lambda^{(q,t)} := E_{\lambda_1} E_{\lambda_2} \cdots E_{\lambda_\ell}(1)$. By Lemma 3.3, $E_k(1) = e_k$, so this is $e_{\lambda_1} \star (e_{\lambda_2} \star (\cdots \star e_{\lambda_\ell}))$. The proof below works for the operators applied in **any** order, so neither associativity nor commutativity of $\star$ is used.

Theorem (DS). For every partition $\lambda \vdash n$, and in Λ_m for every $m \geq n$ (hence in Λ , by §5):

$$e_\lambda^{(q,t)} = s^{n(\lambda)} e_\lambda + \sum_{\mu \triangleright \lambda} c_{\lambda\mu} e_\mu, \quad n(\lambda) = \sum (i-1)\lambda_i,$$

with every $c_{\lambda\mu} \in \mathbb{Q}(t)[s]$ and $c_{\lambda\mu}|_{s=1} = 0$ for $\mu \neq \lambda$. Its parts:

- **(S) Support:** only $\mu \supseteq \lambda$ occur.
- **(L) Leading term:** the coefficient of e_λ is $q^{-n(\lambda)}$.
- **(V) Vanishing:** each off-diagonal coefficient is polynomial in s , so in particular regular at $s = 1$, and it vanishes there.

Operator form (Op-DS). For every $k \geq 1$ and every partition μ :

$$e_k \star e_\mu \in s^{\sum_i \min(\mu_i, k)} e_{\mu \cup k} + \text{span}\{e_\nu : \nu \triangleright \mu \cup k\}.$$

This is the “74 operator-level cases” of Day 211, now for all k and μ .

1. From e-dominance to monomial degree bounds

Lemma 1.1. *Let $m \geq n$ and $\rho \vdash n$. Then $\text{span}\{e_\nu : \nu \supseteq \rho'\} = \text{span}\{m_\kappa : \kappa \leq \rho\}$ in Λ_m^n .*

Proof.

1. By Macdonald I (6.6), $e_\nu = \sum_\kappa M_{\nu\kappa} m_\kappa$, where $M_{\nu\kappa}$ counts 0-1 matrices with row sums ν and column sums κ .
2. If $M_{\nu\kappa} \neq 0$, then $\kappa \leq \nu'$. The first j columns of such a matrix hold at most $\sum_i \min(\nu_i, j) = \nu'_1 + \dots + \nu'_j$ ones.
3. Also $M_{\nu\nu'} = 1$, which is (6.7).
4. By Macdonald I (1.11), $\nu \supseteq \rho' \iff \nu' \leq \rho$. So every e_ν on the left lies in the right span.
5. The e_ν ($\nu \vdash n$) are linearly independent in Λ_m , because $m \geq n$. The map $\nu \mapsto \nu'$ is a bijection between the two index sets, so the dimensions agree. $\square$

Degree. For $S \subseteq [m]$ and a nonzero $R \in K(x)$, let $\deg_S R$ be the degree in u of $R|_{x_i \mapsto ux_i \ (i \in S)}$. Here the u -degree of a rational function over $K(x)$ means $\deg \text{num} - \deg \text{den}$.

Properties:

- $\deg_S(R_1 R_2) = \deg_S R_1 + \deg_S R_2$.
- $\deg_S(R_1 + R_2) \leq \max(\deg_S R_1, \deg_S R_2)$.
- For a polynomial P , $\deg_S P$ is the largest S -degree $\sum_{i \in S} \alpha_i$ of a monomial x^α of P . The reason is that the u^d coefficient is the S -degree- d part of P , which is nonzero.
- For symmetric P , $\deg_S P$ depends only on $j = |S|$. Write $D_j(P)$ for it.

Lemma 1.2. *Let $P \in \Lambda_m$ be homogeneous of degree $n = |\rho|$. Then $P \in \text{span}\{m_\kappa : \kappa \leq \rho\}$ if and only if $D_j(P) \leq \rho_1 + \dots + \rho_j$ for all j .*

Proof. The m_κ have disjoint monomial supports, and $D_j(m_\kappa) = \kappa_1 + \dots + \kappa_j$. $\square$

2. The degree lemma $\Rightarrow$ (S)

Lemma 2.1 (degree step). *Let F be a symmetric polynomial and $k \leq m$. For every j ,*

$$D_j(E_k F) \leq D_j(F) + \min(k, j).$$

Proof. Fix S with $|S| = j$ and bound each subset term $R_A = \prod_A^\times \cdot X_A \cdot F(X_{A^c}, sX_A)$ of (0.1).

- **The kernel.** $\deg_S a_{ij} = 0$ in all cases:
 - if i and j are both in S , or both out, then a_{ij} is homogeneous of degree 0 in u , or constant;
 - if exactly one is in S , then numerator and denominator both have u -degree 1.
- **The monomial.** $\deg_S X_A = |A \cap S| \leq \min(k, j)$.

- **The shifted F .** $F(X_{A^c}, sX_A)$ is F with some variables multiplied by the constant $s \in K^\times$, and this does not change u -degrees. So its $\deg_S$ equals the maximal S -degree of a monomial of F , which is $D_j(F)$.

Hence $\deg_S R_A \leq D_j(F) + \min(k, j)$. The degree of a sum is at most the maximum of the degrees. $E_k F$ is a polynomial by 207b, so its $\deg_S$ is the monomial degree. $\square$

Proof of (S). Induct on the operators applied, in any order. Start from $D_j(1) = 0$. After applying $E_{\lambda_\ell}, \dots, E_{\lambda_1}$,

$$D_j(e_\lambda^{(q,t)}) \leq \sum_i \min(\lambda_i, j) = \lambda'_1 + \dots + \lambda'_j.$$

By Lemma 1.2, $e_\lambda^{(q,t)} \in \text{span}\{m_\kappa : \kappa \leq \lambda'\}$. By Lemma 1.1 this is $\text{span}\{e_\nu : \nu \supseteq \lambda\}$. $\blacksquare$ (S)

Op-DS support is the same argument with a single step: starting from $F = e_\mu$ with $D_j = \sum \min(\mu_i, j)$, one gets $D_j \leq \sum \min(\mu_i, j) + \min(k, j) =$ the partial sums of $(\mu \cup k)'$.

3. The leading coefficient (L)

Fix integers $w_1 > w_2 > \dots > w_m > 0$. Substitute $x_i \mapsto u^{w_i} x_i$ and expand in $K(x)((u^{-1}))$. For $\rho = (\rho_1 \geq \dots \geq \rho_m \geq 0)$, put $W(\rho) := \sum w_i \rho_i$.

Lemma 3.1. *Let x^α be a monomial with $\text{sort}(\alpha) = \kappa \leq \rho$. Then $\sum w_i \alpha_i \leq W(\rho)$, with equality only if $\alpha = \rho$.*

Proof.

- $\sum w_i \alpha_i \leq \sum w_i \kappa_i$ by the rearrangement inequality. Since the w_i are strictly decreasing, equality holds only if α is weakly decreasing, i.e. $\alpha = \kappa$.
- Abel summation gives $\sum w_i \kappa_i = \sum_j (w_j - w_{j+1})(\kappa_1 + \dots + \kappa_j)$, with $w_{m+1} := 0$. Every weight $w_j - w_{j+1}$ is positive, so this is at most the same expression for ρ . Equality holds only if all the partial sums agree, i.e. $\kappa = \rho$. $\square$

Consequently, if $P \in \text{span}\{m_\kappa : \kappa \leq \rho\}$ has m_ρ -coefficient c , then the $u^{W(\rho)}$ -coefficient of $P(u^w x)$ is $c \cdot x^\rho$.

Lemma 3.2 (lead step). *Let $F \in \text{span}\{m_\kappa : \kappa \leq \rho\}$ with m_ρ -coefficient c_F , and let $k \leq m$. Then $E_k F \in \text{span}\{m_\kappa : \kappa \leq \rho + 1^k\}$, and its $m_{\rho+1^k}$ -coefficient is*

$$s^{\rho_1 + \dots + \rho_k} c_F.$$

Proof.

1. **Membership.** This is Lemma 2.1 together with Lemma 1.2. The partial sums of $\rho + 1^k$ are those of ρ plus $\min(k, j)$.
2. **Expand each factor of R_A in u^{-1} .**
 - **The kernel.** If $w_i > w_j$ ($i < j$), then $a_{ij} = (1 - tu^{w_j - w_i} x_j / x_i) / (1 - u^{w_j - w_i} x_j / x_i) = 1 + O(u^{-1})$. If $i > j$, then $a_{ij} = t + O(u^{-1})$. So $\prod_A^\times = t^{\text{inv}(A)} (1 + O(u^{-1}))$, where $\text{inv}(A) = \#\{i \in A, j \notin A : i > j\}$.
 - **The monomial.** X_A becomes $u^{\sum_{a \in A} w_a} X_A$.

- **The shifted F .** $F(X_{A^c}, sX_A)$ has the same monomials as F , but x^α is rescaled by $s^{\sum_{a \in A} \alpha_a}$. By Lemma 3.1 it is $u^{W(\rho)}(c_F s^{\sum_{a \in A} \rho_a} x^\rho + O(u^{-1}))$.

3. **Leading orders.** So $R_A = u^{W(\rho) + \sum_A w_a} (t^{\text{inv}(A)} c_F s^{\sum_A \rho_a} x^{\rho+1_A} + O(u^{-1}))$.

- $\sum_{a \in A} w_a$ is uniquely maximised at $A = [k]$, where $\text{inv} = 0$.
- Hence the $u^{W(\rho+1^k)}$ -coefficient of $E_k F(u^w x)$ is $c_F s^{\rho_1 + \dots + \rho_k} x^{\rho+1^k}$. The other A are of strictly lower order.

4. **Conclude.** Apply the Consequence of Lemma 3.1 to $E_k F$ with the partition $\rho + 1^k$. $\square$

Proof of (L). Apply the operators in any order $p_1, p_2, \dots, p_\ell$, one after another. When E_{λ_p} is applied to the current $\rho = (\cup_i \text{already applied } \lambda_i)'$, it multiplies the lead by

$$s^{\rho_1 + \dots + \rho_{\lambda_p}} = s^{\sum_i \text{applied } \min(\lambda_i, \lambda_p)}.$$

Over the whole process every unordered pair $\{i, p\}$ contributes $\min(\lambda_i, \lambda_p)$ exactly once. So the total exponent is

$$\sum_{i < p} \min(\lambda_i, \lambda_p) = \sum_{i < p} \lambda_p = \sum_p (p-1) \lambda_p = n(\lambda).$$

The $m_{\lambda'}$ -coefficient equals the e_λ -coefficient: in Lemma 1.1's triangularity, the only e_ν ($\nu \succeq \lambda$) that contains $m_{\lambda'}$ is $\nu = \lambda$, and there with coefficient 1. So $c_{\lambda\lambda} = s^{n(\lambda)} = q^{-n(\lambda)}$. $\blacksquare$ (L)

For Op-DS, the lead is $s^{\sum_{i \leq k} \mu'_i} = s^{\sum_i \min(\mu_i, k)}$.

4. Vanishing at $s = 1$ (V)

Let $R := \mathbb{Q}(t)[s]$ and $V := \prod_{i < j} (x_i - x_j)$.

1. **Clearing denominators.** Multiplying (0.1) by V gives

$$V \cdot E_k F = \sum_A \varepsilon_A \cdot \prod_{\substack{i < j \\ \text{both in } A \text{ or both out}}} (x_i - x_j) \cdot \prod_{i \in A, j \notin A} (x_i - tx_j) \cdot X_A \cdot F(X_{A^c}, sX_A),$$

where $\varepsilon_A = (-1)^{\text{inv}(A)}$. Call this sum N . If $F \in R[x]$, then $N \in R[x]$.

2. **Coefficients stay in R .** $V \in \mathbb{Z}[x]$ has content 1. By Gauss's lemma over the UFD $R[x]$, $E_k F = N/V$ lies in $R[x]$. By induction, $e_\lambda^{(q,t)} \in R[x]^{S_m}$. Expanding a symmetric polynomial over a ring R in the e -basis keeps the coefficients in R , so every $c_{\lambda\mu} \in \mathbb{Q}(t)[s]$.
3. **Specialising.** Let $\text{ev} : R[x] \rightarrow \mathbb{Q}(t)[x]$ be the ring map $s \mapsto 1$. Applying it to N kills every s inside $F(X_{A^c}, sX_A)$, so

$$V \cdot \text{ev}(E_k F) = \text{ev}(F) \cdot \sum_A \varepsilon_A(\dots) = \text{ev}(F) \cdot V \cdot E_k(1) = V \cdot \text{ev}(F) \cdot e_k.$$

Cancelling V in the domain $\mathbb{Q}(t)[x]$ gives $\text{ev}(E_k F) = e_k \cdot \text{ev}(F)$.

4. **Iterating.** $\text{ev}(e_\lambda^{(q,t)}) = e_\lambda$. Since the e -expansion is unique, $c_{\lambda\lambda}|_{s=1} = 1$ (consistent with $s^{n(\lambda)}$) and $c_{\lambda\mu}|_{s=1} = 0$ for $\mu \neq \lambda$. $\blacksquare$ (V)

Remark. At $q = 1$ the $\star$ -product is literally the ordinary product. For length 2, (V) was Lemma V ($F_n|_{s=1} = 0$). Lemma V is now a corollary of this structural fact, not a coincidence of the F_n .

5. Stability in m

Let $F \in \Lambda_m$ and set $x_m = 0$ in (0.1). The terms are rational, but their denominators $x_i - x_m$ become $x_i \neq 0$, so the substitution is legitimate.

- **Terms with $m \in A$** die, because of the factor X_A .
- **Terms with $m \notin A$** have $a_{im}|_{x_m=0} = 1$, and $F(X_{A^c}, sX_A)|_{x_m=0} = (F|_{x_m=0})(\dots)$.

So $E_k^{(m)}F|_{x_m=0} = E_k^{(m-1)}(F|_{x_m=0})$. The restriction $\Lambda_m \rightarrow \Lambda_{m-1}$ sends $e_\nu \mapsto e_\nu$, and it is injective in degree n when $m-1 \geq n$. Hence the $c_{\lambda\mu}$ do not depend on $m \geq n$, and DS holds in $\Lambda \otimes \mathbb{Q}(q, t)$. ■

DS is proved for all λ . ■

6. Verification (scripts/day214/)

Engine. `ek_subset_engine.py` implements (0.1) exactly, as N/V with the division checked exact, and e -expands the result.

check_vs_207b.py. Compares (0.1) with the 207b Pieri closed form, *symbolically* in s, t , for $k \leq 3$ and $r \leq 3$ (log `check_vs_207b.log`). This confirms the conventions and normalisation match the proved length-2 theory.

check_ds_all_lengths.py N. Runs over every $\lambda \vdash n \leq N$, of all lengths, at the exact point $(s, t) = (3/7, -5/11)$, and checks:

- (S);
- (L), i.e. $\text{lead} = s^{n(\lambda)}$ exactly;
- (V), by running the engine at $s = 1$ and checking $e_\lambda^{(q,t)} = e_\lambda$;
- order-independence, $E_{\lambda_\ell} \cdots E_{\lambda_1} = E_{\lambda_1} \cdots E_{\lambda_\ell}$;
- “support = full up-set”.

Results:

- $N = 5$: all 18 partitions pass everything (`check_ds_all_lengths_N5.log`).
- $N = 6$: see `check_ds_all_lengths_N6.log`.

How to read the full-up-set results. The $c_{\lambda\mu}$ are polynomials in s with coefficients in $\mathbb{Q}(t)$. So a nonzero value at an exact rational point *proves* $c_{\lambda\mu} \neq 0$. For the partitions checked, “support = full up-set” is therefore proved case by case.

Independent confirmation. Day 211’s (TC)-based engine found full up-set support for length 3, $|\lambda| \leq 6$.

7. Gaps, caveats, scope

- **No gap in the argument.**
 - The load-bearing inputs are 207b (A_k) and (K_k), both PROVED, reviewed by Clio on 2026-09-29, together with Lemma 3.3, a Hikita verified-quote (R0).
 - Macdonald I (1.11), (6.6) and (6.7) are textbook.
 - Gauss’s lemma is textbook.
- **Stretch goal: PROVED in §8.** Every $c_{\lambda\mu}$ with $\mu \supseteq \lambda$ is nonzero, and in fact $\text{val}_s c_{\lambda\mu} = n(\mu)$ exactly.
- **Trust.** DS no longer depends on (TC) or ($\star\ell$). So Clio’s pending review of those does NOT gate DS. DS inherits only the 207b grade.
- **Novelty caution.** The argument is the standard “Macdonald-operator triangularity” degree count (compare Macdonald VI §3 / VI (3.6)–(3.10): D_n^r acts triangularly on m_λ). Since the $\star$ -operators have exactly Macdonald’s shape — a kernel $\prod^\times$ times a shift, here twisted by the X_A factor — it is plausible that Hikita or someone else states DS, or its monomial-triangular form, explicitly. That needs a novelty check before any claim. No browsing was done in this session.
- **What changed about the Day 211 obstruction.** The (TC) termwise dominance violations are real, but they are an artefact of the $E(t^i z)E(t^j w)$ expansion basis. The subset expansion (0.1) is termwise degree-bounded. **Lesson: choose the expansion in which the bound is termwise.**

8. Stretch goal PROVED: the support is exactly the up-set, and $\text{val}_s c_{\lambda\mu} = n(\mu)$

Theorem 2. *Let λ be a partition and $\mu \supseteq \lambda$.*

- (a) $c_{\lambda\mu} \in \mathbb{Q}[s, t]$, and $s^{n(\mu)}$ divides $c_{\lambda\mu}$.
- (b) $c_{\lambda\mu}(s, 0) = s^{n(\mu)}(1 + s \cdot \mathbb{Q}[s])$.

Consequently:

- every $c_{\lambda\mu}$ ($\mu \supseteq \lambda$) is nonzero, so $\text{supp } e_\lambda^{(q,t)} = \{\mu : \mu \supseteq \lambda\}$ exactly;
- the s -adic valuation of $c_{\lambda\mu}$ is exactly $n(\mu)$;
- the lowest coefficient $d_{\lambda\mu}(t) := [s^{n(\mu)}]c_{\lambda\mu} \in \mathbb{Q}[t]$ satisfies $d_{\lambda\mu}(0) = 1$.

At $s = t = 0$, in the rescaled basis $b_\mu := s^{n(\mu)}e_\mu$, $e_\lambda^{(q,t)}$ is $\sum_{\mu \supseteq \lambda} b_\mu$, **the zeta function of the dominance order.**

Notation.

- $\omega(\alpha) := \sum_j \binom{\alpha_j}{2}$ for $\alpha \in \mathbb{Z}_{\geq 0}^m$.
- $B_\mu := \{\beta \in \mathbb{Z}_{\geq 0}^m : \text{sort}(\beta) \leq \mu'\}$. By §2, this is the monomial support set of $e_\mu^{(q,t)}$.

- Write $e_\mu^{(q,t)} = \sum_\beta a_\beta^\mu x^\beta$. The coefficient a_β^μ is symmetric in β , because $e_\mu^{(q,t)}$ is a symmetric polynomial.

Key identity. For $A \subseteq [m]$, $\omega(\beta + 1_A) = \omega(\beta) + 1_A \cdot \beta$, since $\binom{b+1}{2} = \binom{b}{2} + b$. This is where s^n comes from.

8.1 Integrality

§4's Gauss-lemma argument works over $\mathbb{Q}[s, t]$, because N_A uses only $(x_i - tx_j)$ and no t^{-1} . So E_k preserves $\mathbb{Q}[s, t][x]^{S_m}$, and hence $c_{\lambda\mu} \in \mathbb{Q}[s, t]$.

Specialising $t = 0$ is a ring map. Moreover $c_{\lambda\mu}(s, 0)$ are the e -coefficients of $E_{\lambda_1}^0 \cdots E_{\lambda_\ell}^0(1)$, where E_k^0 is (0.1) with a_{ij} replaced by $x_i/(x_i - x_j)$.

8.2 Lemma A (valuation bound)

Lemma A (valuation bound). For t generic, and also at $t = 0$, $\text{val}_s(a_\alpha^\lambda) \geq \omega(\alpha)$ for all α .

Setup.

- Work over $K = \mathbb{Q}(t)(s^{1/2})$, or $K = \mathbb{Q}(s^{1/2})$ at $t = 0$, with the s -adic valuation v . It is trivial on $\mathbb{Q}(t)$.
- Extend v to $K(x)$ by the Gauss valuation: $v(\sum a_\beta x^\beta) = \min_\beta v(a_\beta)$, and $v(P/Q) = v(P) - v(Q)$. This is a valuation.
- For $c \in (\frac{1}{2}\mathbb{Z})^m$, let σ_c be the K -automorphism $x_j \mapsto s^{-c_j} x_j$.

Claim. $v(\sigma_c e_\lambda^{(q,t)}) \geq \min_{\beta \in B_\lambda} (\omega(\beta) - c \cdot \beta)$ for all c .

Proof of the claim. Induct along E_k : $e_\mu \mapsto e_{\mu \cup k}$. Base case: $1 = e_\emptyset$, with $B_\emptyset = \{0\}$.

For a term R_A of (0.1) applied to $F = e_\mu^{(q,t)}$, the three factors of $\sigma_c R_A$ are:

- **The kernel.** At generic t , $v(\sigma_c a_{ij}) = \min(-c_i, -c_j) - \min(-c_i, -c_j) = 0$. At $t = 0$, $v = -c_i + \max(c_i, c_j) \geq 0$.
- **The monomial.** $\sigma_c X_A = s^{-c \cdot 1_A} X_A$.
- **The shifted F .** $\sigma_c[F(X_{A^c}, sX_A)] = \sigma_{c-1_A} F$.

Hence

$$v(\sigma_c R_A) \geq -c \cdot 1_A + \min_{\beta \in B_\mu} (\omega(\beta) - (c - 1_A) \cdot \beta) = \min_{\beta \in B_\mu} (\omega(\beta + 1_A) - c \cdot (\beta + 1_A)),$$

using the key identity. Now $\beta + 1_A \in B_{\mu \cup k}$: the top- j sums grow by at most $\min(j, k)$, and $\mu' + 1^k = (\mu \cup k)'$. So the bound is $\geq \min_{B_{\mu \cup k}}$. Take the minimum over A . $\square$

Extraction. Put $c := \alpha - \frac{1}{2} \cdot \mathbb{1}$. Then

$$\omega(\beta) - c \cdot \beta = \sum_j \left[\frac{(\beta_j - \alpha_j)^2}{2} - \frac{\alpha_j^2}{2} \right],$$

whose *unique* minimiser over all of $\mathbb{Z}^m$ is $\beta = \alpha$. Since $v(\sigma_c P) = \min_\beta (v(a_\beta) - c \cdot \beta)$, we get $v(a_\alpha) - c \cdot \alpha \geq \omega(\alpha) - c \cdot \alpha$. $\blacksquare(A)$

8.3 From monomials to e -coefficients

Recall $n(\nu) = \sum_i (n - S_i(\nu))$, so $\nu \triangleleft \mu$ implies $n(\nu) > n(\mu)$. Also $n(\mu) = \omega(\mu')$.

Write $c_\nu = s^{n(\nu)} d_\nu$. Suppose $\min_\nu v(d_\nu) = \delta < 0$. Among the ν with $v(d_\nu) = \delta$, pick μ minimal in dominance. Then:

- $a_{\mu'} = \sum_{\nu \trianglelefteq \mu} c_\nu M_{\nu\mu'}$, with $M_{\mu\mu'} = 1$ (Lemma 1.1).
- The $\nu = \mu$ term has valuation exactly $n(\mu) + \delta$.
- Every other term has valuation $> n(\mu) + \delta$: if $\nu \triangleleft \mu$, then $n(\nu) > n(\mu)$.

So $v(a_{\mu'}) = n(\mu) + \delta < \omega(\mu')$, contradicting Lemma A. Hence every d_ν is s -integral, which proves Theorem 2(a). Reducing mod s gives

$$d_\mu(0) = [s^{\omega(\mu')}] a_{\mu'} =: \Lambda_\lambda(\mu').$$

8.4 Lemma B (initial forms at $t = 0$)

Lemma B (initial forms at $t = 0$). At $t = 0$, $\Lambda_\lambda^0(\alpha) := [s^{\omega(\alpha)}] a_\alpha^\lambda$ equals 1 if $\text{sort}(\alpha) \leq \lambda'$, and 0 otherwise.

Setup. Let $\text{in}_{v_0}(R)$ denote the reduction of $s^{-v_0}R$ in the residue field $\mathbb{Q}(x)$ of the Gauss valuation, defined when $v(R) \geq v_0$. It is additive at a common level, and multiplicative.

Fix α with $\text{sort } \alpha \leq (\mu \cup k)'$, put $c = \alpha - \frac{1}{2}$ and $v_0 = \omega(\alpha) - c \cdot \alpha$. By the extraction uniqueness,

$$\text{in}_{v_0}(\sigma_c e_{\mu \cup k}^{(q,t)}) = \Lambda_{\mu \cup k}^0(\alpha) x^\alpha.$$

Termwise analysis. Consider the term R_A .

- **Terms that miss the top order.** Suppose $\alpha - 1_A$ is not in B_μ , or has a negative entry. Then $v(\sigma_c R_A) > v_0$ strictly, because the unique minimiser α is not of the form $\beta + 1_A$. So in_{v_0} kills R_A .
- **Terms that survive.** Otherwise, with $\beta_0 = \alpha - 1_A$ and $c - 1_A = \beta_0 - \frac{1}{2}$,

$$\text{in}_{v_0}(\sigma_c R_A) = \text{in}_0(\sigma_c \prod_A^\times) \cdot \Lambda_\mu^0(\beta_0) \cdot x^\alpha.$$

The kernel factor at $t = 0$. Here $\text{in}_0(\sigma_c a_{ij})$ equals:

- 1 if $\alpha_i > \alpha_j$;
- 0 if $\alpha_i < \alpha_j$;
- $x_i/(x_i - x_j)$ if $\alpha_i = \alpha_j$.

So only A that are *upper sets* for α survive. Let v be the k -th largest entry of α and $L_v := \{j : \alpha_j = v\}$. The surviving sets are $A = \{\alpha > v\} \sqcup B$, where $B \subseteq L_v$ has size $r := k - \#\{\alpha > v\}$. Every entry of A is ≥ 1 , because $\ell(\text{sort } \alpha) \geq \ell((\mu \cup k)') \geq k$.

Summing over B . $\text{sort}(\alpha - 1_A) = \text{peel}_k(\text{sort } \alpha)$ does not depend on B , and Λ_μ^0 is symmetric. So

$$\Lambda_{\mu \cup k}^0(\alpha) = \Lambda_\mu^0(\text{peel}_k \alpha) \cdot \sum_{B \subseteq L_v, |B|=r} \prod_{i \in B, j \in L_v \setminus B} \frac{x_i}{x_i - x_j} = \Lambda_\mu^0(\text{peel}_k \alpha).$$

The last sum is 1. It is symmetric (S_{L_v} permutes the terms), so its product with the Vandermonde is antisymmetric, and the sum is a polynomial. It is homogeneous of degree 0, hence a constant. The constant is its u -leading coefficient under the weights $w_1 > w_2 > \dots$ of §3, where only B = the first r indices contributes, with value 1. The t -version, $\Sigma = \binom{N}{r}_t$, is checked in `check_stretch.py`.

Here $\text{peel}_k(\kappa)$ means: subtract 1 from the k largest parts of κ .

Peel Lemma. *Let $\kappa \trianglelefteq \lambda$ with $|\kappa| = |\lambda|$ and $k \leq \ell(\lambda)$. Then $\text{peel}_k \kappa \trianglelefteq \text{peel}_k \lambda$.*

Proof.

1. **Conjugate formula.** Counting the parts $\geq c$ after the peel gives

$$(\text{peel}_k \kappa)'_c = \kappa'_c - \min(k, \kappa'_c) + \min(k, \kappa'_{c+1}).$$

(This is checked for $n \leq 12$ in `check_stretch.py`.)

2. **Partial sums.** Telescoping, and using $\kappa'_1 = \ell(\kappa) \geq \ell(\lambda) \geq k$,

$$S_C((\text{peel}_k \kappa)') = S_C(\kappa') - k + \min(k, \kappa'_{C+1}).$$

The same holds for λ .

3. **Compare.** We have $\kappa' \trianglerighteq \lambda'$ (Macdonald I (1.11)). Put $D_C := S_C(\kappa') - S_C(\lambda') \geq 0$ and $m(x) := \min(k, x)$.

- If $m(\kappa'_{C+1}) \geq m(\lambda'_{C+1})$, we are done.
- Otherwise $\kappa'_{C+1} < k$ and $\kappa'_{C+1} < \lambda'_{C+1}$. Then $D_C \geq D_{C+1} + \lambda'_{C+1} - \kappa'_{C+1} \geq m(\lambda'_{C+1}) - m(\kappa'_{C+1})$.

4. So $(\text{peel}_k \kappa)' \trianglerighteq (\text{peel}_k \lambda)'$, which is $\text{peel}_k \kappa \trianglelefteq \text{peel}_k \lambda$. □

(A brute-force check confirms it for all 10,788 triples with $n \leq 11$, in `peel_monotone.py`. This is the greedy step of Ryser's algorithm for the Gale–Ryser theorem.)

Proof of Lemma B. Induct along E_k . The base case is $\Lambda_{\emptyset}^0(0) = 1$. For sort $\alpha \trianglelefteq (\mu \cup k)'$, the Peel Lemma with $\lambda := (\mu \cup k)'$ gives:

- $\text{peel}_k(\text{sort } \alpha) \trianglelefteq \text{peel}_k((\mu \cup k)')$;
- $\text{peel}_k((\mu \cup k)') = \mu'$, because $(\mu \cup k)' = \mu' + 1^k$ and its top k parts sit in positions $1..k$.

So $\Lambda_{\mu \cup k}^0(\alpha) = 1$. Outside $B_{\mu \cup k}$, the value is 0 by (S). ■(B)

Proof of Theorem 2. $d_{\lambda\mu}(0)|_{t=0} = \Lambda_{\lambda}^0(\mu') = [\mu' \trianglelefteq \lambda'] = [\mu \trianglerighteq \lambda] = 1$. Since $c_{\lambda\mu}/s^{n(\mu)} \in \mathbb{Q}[s, t]$ (by §8.3, together with $s^{n(\mu)} \mid c$ in $\mathbb{Q}(t)[s] \cap \mathbb{Q}[s, t]$), its value at $s = t = 0$ is 1. ■

Checks.

- `s_valuation.py` ($n \leq 5$), `check_stretch.py` ($n \leq 5$) and `check_stretch_n6.py` ($n = 6$, $t = 0$): $\text{val}_s c_{\lambda\mu} = n(\mu)$ for every $\mu \trianglerighteq \lambda$, with lowest coefficient exactly 1 at $t = 0$. At $t = -5/11$ the valuation is again $n(\mu)$.
- The $d_{\lambda\mu}(t)$ are genuinely t -dependent. Example: $d_{(1111), (211)} = 1 + 3t$, a value consistent with the data at $t = -5/11$.

- *Remark (computed for $n \leq 5$ in `t1_count.py`; the $t = 1$ case follows from the same in_0 argument, but it is not written out).* At $t = 1$ the recursion becomes $\Lambda_{\mu \cup k}(\alpha) = \sum_{A \subseteq \text{supp } \alpha} \Lambda_{\mu}(\alpha - 1_A)$. This counts 0-1 matrices, so $d_{\lambda\mu}(1) = M_{\lambda\mu'}$. The case $1+3t$ at $t = 1$ gives $4 = M_{(1111), (31)}$. At general t , the in_0 of a_{ij} for $\alpha_i < \alpha_j$ is t , and the level-set sums are t -binomials. So $d_{\lambda\mu}(t)$ is a t -count of 0-1 matrices with row sums λ and column sums μ' . This is not written out here.